

init rides the rocket: systemd is here

Olaf Kirch

Director SUSE Linux Enterprise

okir@suse.com



Love it or hate it?



1996: Linux Distros adopt SysV-init



2001: LSB standardizes init scripts



2010: Mobile Distros and Parallel Boot



2027: Still in Good Shape?





HOW ABOUT
ENDING THIS
CARTOON
BEFORE I
HIT?

Why did you do this to me?!

What's wrong with sysvinit?

- Nothing
- But it could do many things better
 - It's slow
 - It's hard to parallelize
 - Too coarse synchronization points
 - LSB dependencies only do what you need 50% of the time
 - Have you ever tried to kill all processes spawned by a user session?
 - No automatic restart of services
 - No unified logging
 - No unified resource limit handling

Why systemd

- Considered several alternatives, systemd came out on top for several reasons
- It's not Marmite, but it certainly polarizes the users today
 - Some of that is certainly because it's new
- It seems to be the solution most distros are standardizing on

Why is it so hard to like systemd, then?

- It isn't, once you get to know to it better
- It changes a lot of things
 - You need to learn a lot of new commands
 - But let's be honest, many of them do things you could never do with sysvinit
 - Some time-honored features (like `inittab`) are simply gone
 - It is an intrusive change, and fixing up the fallout has been a significant amount of work
 - “Backward compatibility” is not high on the list of development priorities
- That said, some level of backward compatibility is possible, and we're providing that

SUSE backward compatibility

- insserv, chkconfig and /sbin/service will still be supported
- Old style “rcfooobar start” redirected to new systemd tools automatically
- LSB compatibility for targets like `$network...` still available
- And of course init scripts are still supported!

Systemd Concepts

Generic concepts

- systemd replaces the traditional init process
 - It's not one, but a collection of DBus services
- Activate services on demand rather than up-front
- One-stop shopping for boot, shutdown and power management
 - integrated logging
 - unified command line tools for almost everything
 - automatic restart of services
 - cgroups and security compartments for everything
 - multi-seat hosts
 - handling of ACPI power management events

Unit files

- Unit files describe *targets* (i.e. synchronization points) and *services* (what used to be init scripts)
- Runlevels are replaced by targets
 - Runlevel 3: multi-user.target
 - Runlevel 5: graphical.target
- Much of what the LSB standard used is modeled in unit files
 - Plus a few more, for instance with LSB you could never say “my service needs to be started before kdm”

Sessions and Seats

- A *seat* is the set of hardware available at one work place (graphics card, keyboard, mouse, usb devices)
- A *session* is created once a user is logged on, using a specific seat
 - Only one session can be active per seat
 - Default seat (for Linux consoles) is **seat0**
- Hardware is assigned to seats
 - This replaces ConsoleKit

Sample service unit file: sshd.service

```
# This is a comment!
```

```
[Unit]
```

```
Description=OpenSSH Daemon
```

```
After=network.target
```

```
[Service]
```

```
EnvironmentFile=/etc/sysconfig/ssh
```

```
ExecStartPre=/usr/sbin/sshd-gen-keys-start
```

```
ExecStart=/usr/sbin/sshd -D $SSHD_OPTS
```

```
ExecReload=/bin/kill -HUP $MAINPID
```

```
KillMode=process
```

```
Restart=always
```

```
[Install]
```

```
WantedBy=multi-user.target
```

```
LSB Analogs/Equivalents
```

```
[Unit]
```

```
# Description: ...
```

```
# Required-Start: $network
```

```
[Service]
```

```
# <- All of these used to be
```

```
# <- open coded in the init script
```

```
# <-
```

```
# <-
```

```
# <-
```

```
Automated Restart: didn't exist previously
```

```
[Install]
```

```
# Default-Start: 3 5
```



Cgroups for Everything

- Systemd puts each service and each session into a separate cgroup
 - Sessions also get assigned an audit ID matching their cgroup ID
- You can restrict these cgroups in all the way the kernel supports
 - IO bandwidth, memory or CPU consumption, etc

Improved Security for Everything

- Restrict services and sessions using namespaces
 - Linux kernel namespaces are the technology underlying Linux containers
 - blacklist directories
 - require private /tmp directory
 - whitelist devices to which access is granted
- specify user/group to run as
- assign Linux kernel capabilities (CAP_FOOBAR)
- set ulimit values

Overriding defaults for a service

- With sysvinit, if you want to do anything more advanced that enable/disable a service, you need to edit the init script
 - This doesn't go well with security updates
- Systemd supports that nicely
 - Assume you want to modify settings for **foobar.service**
 - Create **/etc/systemd/system/foobar.service.d**
 - Drop a file named **mysettings.conf** in there:

```
[Service]
```

```
InaccessibleDirectories=/precious
```

```
MemoryLimit=1G
```

Getting Started with Systemd Tools

Very Rough Cheat Sheet

- The following overview is far from exhaustive
- This is just meant as a starting point to help you exploring systemd and its tools

Service status

- You want a list of all started services and their status

```
systemctl
```

- You want the status of service **foobar**:

```
systemctl status foobar.service
```

```
$ systemctl status icecream.service
icecream.service - LSB: icecc
   Loaded: loaded (/etc/init.d/icecream)
   Active: active (running) since Fri, 2013-04-19 09:27:31 CEST; 4 days ago
   CGroup: name=systemd:/system/icecream.service
           └─ 4786 /usr/sbin/icecc-scheduler -d -l /var/log/icecc_sch...
              4791 /usr/sbin/iceccd -d -l /var/log/iceccd --nice 5 -u...

Apr 19 09:27:31 foobar systemd[1]: Starting LSB: icecc...
Apr 19 09:27:31 foobar icecream[4777]: Starting Distribut...
Apr 19 09:27:31 foobar systemd[1]: Started LSB: icecc.
```


Starting and Stopping Services

- `systemctl <verb> foobar.service`
 - Where *<verb>* is one of start, stop, restart, try-restart, reload
- `systemctl kill foobar.service`
 - Kill all processes in the cgroup of this service

Enabling and Disabling Services

- `systemctl <verb> foobar.service`
 - Where *<verb>* is one of enable, disable

We'll still be friends, pstree

```
# systemd-cgls
```

```
system
├─ 1 /sbin/init showopts
├─ icecream.service
│   └─ 4786 /usr/sbin/icecc-scheduler -d -l /var/log/icecc_scheduler
│       └─ 4791 /usr/sbin/iceccd -d -l /var/log/iceccd --nice 5 -u icecream -b /...
├─ colord.service
│   └─ 1677 /usr/lib/colord
├─ udisks2.service
│   └─ 1498 /usr/lib/udisks2/udisksd --no-debug
├─ rtkit-daemon.service
│   └─ 1353 /usr/lib/rtkit/rtkit-daemon
├─ upower.service
│   └─ 1161 /usr/lib/upower/upowerd
├─ accounts-daemon.service
│   └─ 1125 /usr/lib/accounts-daemon
├─ xdm.service
│   └─ 964 /usr/sbin/gdm
│       └─ 966 /usr/lib/gdm/gdm-simple-slave --display-id /org/gnome/DisplayMan...
│           └─ 1021 /usr/bin/Xorg :0 -background none -verbose -auth /run/gdm/auth-f...
│               └─ 1515 gdm-session-worker [pam/gdm-password]
```

Session Handling

- List all sessions:

- loginctl [list-sessions]

- Show session details:

- loginctl session-status <session-number>

- Forcefully terminate a session:

- loginctl kill-session|kill-user|terminate-seat <name>

Session Handling, continued

```
2 - fcrozat (1000)
    Since: lun. 2013-07-29 11:58:41 CEST; 4h 13min ago
    Leader: 1550 (gdm-session-wor)
    Seat: seat0; vc7
    Display: :0
    Service: gdm-password; type x11; class user
    State: active
    CGroup: systemd:/user/1000.user/2.session
            1550 gdm-session-worker [pam/gdm-password]
            1557 /usr/bin/gnome-keyring-daemon --daemonize
            1560 /usr/bin/gnome-session
```


And it comes with lots more stuff

- See how your configuration differs from the vendor defaults
 - `systemd-delta`
- Analyze boot times and bottlenecks
 - `systemd-analyze`

References

- Systemd on SLE12 :
 - Check our official documentation
- Upstream:
 - <http://www.freedesktop.org/wiki/Software/systemd/>
 - Check manpages, they are extremely verbose (if something is missing there, it is a bug !)

You are now ready to test
systemd in SLE12

Thank you.

